

01 Cloud PubSub Interview Questions - KirkYagami

Publish-Subscribe (PubSub) Systems

1. What is a publish-subscribe (pubsub) system, and how does it differ from point-to-point messaging?

A publish-subscribe (pubsub) system is a messaging pattern where senders (publishers) send messages to topics without needing to know who the receivers (subscribers) are. Subscribers listen to specific topics and receive all messages published on those topics.

```
- **Publisher**: The news agency                                     language-txt
- **Subscribers**: All users subscribed to the "Sports" topic
- **Message**: New article about a football match
- **Topic**: "Sports"
```

```
# Example service: Google Cloud Pub/Sub, where messages are sent to a topic and
delivered to all subscribers.`
```

In contrast, point-to-point messaging involves direct communication between a sender and a receiver, where each message is consumed by only one receiver.

```
- **Sender**: User A                                                 language-txt
- **Receiver**: User B
- **Message**: A personal email from User A to User B
- **Message Queue**: The email is processed and delivered to only User B
```

```
# Example service: Message Queue systems like Amazon SQS, where each message is
consumed by only one recipient.
```

2. What is Google Cloud Pub/Sub, and what problem does it solve?

Google Cloud Pub/Sub is a fully managed messaging service that allows for asynchronous communication between applications in a scalable, reliable, and decoupled manner. It helps solve the problem of real-time data ingestion, streaming, and inter-service communication in a distributed system. Pub/Sub is designed to enable applications to communicate via

publish/subscribe messaging patterns, where a publisher sends messages to a topic and one or more subscribers receive them asynchronously.

Key problems it solves:

- **Decoupling microservices**: Services do not need to know the details of one another to communicate.
- **Scalability**: Designed to handle millions of messages per second.
- **Reliability**: Ensures messages are delivered at least once and can handle message ordering and deduplication.

[Official Docs: Google Cloud Pub/Sub Overview](#)

3. Explain the main components of a cloud pubsub architecture.

In a **cloud Pub/Sub architecture**, the main components include:

1. Publishers

- **Definition**: Services or applications that send messages to the system.
- **Function**: Publishers don't send messages directly to subscribers but instead send them to topics. Publishers can be a part of any microservice or event-generating system.
- **Example**: A server publishing updates about stock prices, sensor data from IoT devices, or user activities from a web application.

2. Topics

- **Definition**: Named resources to which messages are sent by publishers.
- **Function**: Topics serve as the medium that connects publishers to subscribers. They categorize messages based on their type or subject. Messages from a publisher are published to a specific topic.
- **Example**: In a stock trading system, there could be topics like "stock-price-updates" and "trade-confirmations."

3. Messages

- **Definition**: Data sent by the publisher to the topic.
- **Function**: Each message contains a payload (the actual data) and metadata (such as attributes or labels). The metadata helps route the message appropriately.
- **Example**: A stock price update message might contain the stock symbol, current price, and timestamp.

4. Subscriptions

- **Definition:** A subscription defines the relationship between a topic and a subscriber.
- **Function:** Subscribers register to a specific topic by creating a subscription. A subscription will receive all messages published to its associated topic. Each subscription can have its own unique configurations, like the message delivery method (push or pull).
- **Types of Subscriptions:**
 - **Pull:** The subscriber must explicitly request (or pull) messages from the subscription.
 - **Push:** The system sends (or pushes) messages directly to the subscriber.

5. Subscribers

- **Definition:** Services or applications that consume messages from subscriptions.
- **Function:** Subscribers process the messages they receive from a subscription. Multiple subscribers can listen to the same topic, each with its own subscription.
- **Example:** In a microservices architecture, different services (like billing, analytics, or notifications) might subscribe to the same topic and process the data differently.

6. Message Delivery

- **Definition:** The mechanism by which messages are delivered from topics to subscribers.
- **Function:** Cloud Pub/Sub ensures reliable message delivery by managing retries and ensuring at-least-once delivery. It handles failures or timeouts by retrying until the message is acknowledged by the subscriber.
- **Acknowledgments (ACKs):** Subscribers must acknowledge (ack) that they have received and processed a message successfully. If not, the system may resend the message.

7. Message Retention & Ordering

- **Message Retention:** Messages are stored temporarily if a subscriber is not available. Cloud Pub/Sub can retain unacknowledged messages for a configured time.
- **Message Ordering:** Some systems offer in-order delivery of messages if needed, but most cloud Pub/Sub systems offer "best-effort" ordering by default.

8. Dead Letter Policy (optional)

- **Definition:** A mechanism to handle undeliverable messages.
- **Function:** If a message fails to be delivered after a certain number of attempts, it is moved to a dead letter queue for later inspection, preventing the system from becoming overloaded with undeliverable messages.

9. Monitoring & Scaling

- Cloud Pub/Sub systems often come with built-in monitoring tools to track the number of published messages, the processing rate by subscribers, and errors.
- Pub/Sub systems are designed to scale automatically with demand, handling large volumes of messages across distributed systems without manual intervention.

This architecture decouples producers and consumers, making the system more flexible, scalable, and fault-tolerant.

[Official Docs: Topics and Subscriptions](#)

4. What are the advantages of using a pubsub system in a distributed architecture?

-Using a **Pub/Sub system** in a distributed architecture offers several key advantages:

1. Decoupling of Components

- **Advantage:** Publishers and subscribers are decoupled. Publishers don't need to know about the subscribers or their logic, and vice versa. This allows each component to evolve independently.
- **Benefit:** Systems can scale and adapt more easily since new services or changes to existing ones don't impact other parts of the system.

2. Scalability

- **Advantage:** Pub/Sub systems are designed to handle large volumes of messages and scale horizontally. They can dynamically adjust resources based on the traffic load.
- **Benefit:** You can easily accommodate spikes in traffic, ensuring that the system remains responsive even as the number of publishers, subscribers, or messages grows.

3. Asynchronous Communication

- **Advantage:** Publishers can send messages without waiting for subscribers to process them. Subscribers can process messages when they are ready.
- **Benefit:** This improves system performance and responsiveness, especially in environments with varying workloads or latency-sensitive applications.

4. Reliable Message Delivery

- **Advantage:** Pub/Sub systems provide mechanisms to ensure messages are reliably delivered, often guaranteeing at-least-once delivery.

- **Benefit:** Messages won't get lost even if there are temporary failures in the system. The retry mechanisms ensure fault tolerance and high availability.

5. Real-Time Data Streaming

- **Advantage:** Pub/Sub systems allow for real-time data streams where messages can be processed almost instantaneously after being published.
- **Benefit:** Ideal for scenarios like event-driven systems, sensor data processing, and live user notifications where real-time data flow is crucial.

6. Flexible Scaling of Subscribers

- **Advantage:** Subscribers can process messages independently, allowing more instances of subscribers to be added as needed.
- **Benefit:** This supports dynamic scaling of processing services, ensuring that high traffic or message loads don't overwhelm individual subscribers.

7. Fault Isolation

- **Advantage:** Failures in subscribers do not affect publishers, and vice versa, because of the loose coupling.
- **Benefit:** This enhances system reliability since different parts of the system can fail independently without cascading failures.

8. Load Balancing

- **Advantage:** Pub/Sub systems can distribute messages among multiple subscribers based on availability, ensuring that no single service is overwhelmed.
- **Benefit:** This allows efficient use of resources and helps manage traffic surges.

9. Event-Driven Architecture

- **Advantage:** Pub/Sub systems support event-driven architectures where components react to published events, allowing for more modular, reactive systems.
- **Benefit:** This simplifies designing systems that need to respond to events or triggers, improving modularity and reducing complexity in microservice-based environments.

10. Reduced Latency for Subscribers

- **Advantage:** With Pub/Sub, subscribers can receive messages as soon as they are published without polling or making constant requests.

- **Benefit:** This reduces the overall latency in delivering messages, which is critical for time-sensitive applications like financial markets, IoT systems, or notifications.

11. Simplified Communication in Distributed Systems

- **Advantage:** Pub/Sub abstracts the complexity of point-to-point communication in distributed systems by allowing a many-to-many communication pattern.
- **Benefit:** This reduces the need for hardcoded communication paths and simplifies communication between services.

12. Efficient Resource Utilization

- **Advantage:** Resources are used more efficiently as messages are only processed when subscribers are ready.
- **Benefit:** Prevents wasted CPU cycles or unnecessary network traffic caused by continuous polling, optimizing both performance and costs.

13. Support for Multiple Message Consumers

- **Advantage:** A single message can be delivered to multiple subscribers if necessary.
- **Benefit:** Multiple services can react to the same event, supporting diverse use cases like logging, auditing, and notification systems based on a single published event.

In summary, a **Pub/Sub system** is particularly advantageous in **distributed architectures** due to its ability to **decouple components**, **scale efficiently**, **support asynchronous communication**, and provide **reliable, real-time message delivery**. It simplifies the design and management of large-scale, event-driven systems and improves overall system flexibility and robustness.

5. Can you describe some common use cases for pubsub systems?

- **Event-driven architectures:** Propagating state changes, e.g., in microservices.
- **Real-time data streaming:** Stock market data, sensor networks, or social media feeds.
- **Logging and monitoring:** Collecting logs and metrics across distributed systems.
- **Notification systems:** Sending alerts or updates to multiple recipients.

6. Discuss the difference between synchronous and asynchronous communication in PubSub.

- **Synchronous communication:** The sender waits for a response from the receiver before proceeding.

- **Asynchronous communication:** The sender doesn't wait for a response. PubSub is primarily asynchronous, meaning publishers don't need to wait for subscribers to process messages.

Synchronous Communication in Pub/Sub

In synchronous communication, the publisher and subscriber are more tightly coupled in terms of timing. When the publisher sends a message, it expects an immediate acknowledgment or response.

Key Characteristics:

- **Immediate Acknowledgment:** The publisher sends a message and waits for the subscriber to process the message and send back an acknowledgment before moving on.
- **Blocking Behavior:** The publisher halts further operations until the subscriber confirms receipt and processing of the message.
- **Tighter Coupling:** There is a more direct relationship between the publisher and subscriber, meaning the publisher is dependent on the subscriber's availability and speed of processing.

Example:

- In a **request-response** system, the publisher (requester) sends a message (request) and waits for an immediate response from the subscriber (responder). This might be a web service that sends a request for data and expects an immediate response.

Advantages:

- **Consistency:** The publisher knows immediately whether the message was received and processed, ensuring real-time feedback.
- **Simplicity:** Easier to trace the flow of messages since it's linear—publish, wait, respond.

Drawbacks:

- **Lower Scalability:** If a subscriber is slow or unresponsive, the publisher is forced to wait, creating potential bottlenecks.
 - **Higher Latency:** The system can experience delays since the publisher has to wait for an acknowledgment.
 - **Dependency on Subscriber Availability:** The system can fail if the subscriber is down or unable to process the message in time.
-

Asynchronous Communication in Pub/Sub

In asynchronous communication, the publisher and subscriber operate independently. The publisher sends a message to the topic and does not wait for the subscriber to process it, allowing both to operate concurrently.

Key Characteristics:

- **Non-blocking Behavior:** The publisher sends a message and immediately continues its operations without waiting for any acknowledgment from the subscriber.
- **Decoupled Communication:** The publisher does not depend on the subscriber's state or availability. Subscribers can process messages at their own pace.
- **Message Delivery Guarantees:** While the publisher does not wait for immediate acknowledgment, the system often ensures message delivery through retries or queuing until the subscriber confirms receipt.

Example:

- In an **event-driven architecture**, the publisher might send messages to a topic (e.g., logging user activity, sending notifications), and subscribers consume those messages whenever they are ready to process them.

Advantages:

- **High Scalability:** Publishers and subscribers are decoupled, allowing the system to handle high volumes of messages without blocking or waiting.
- **Fault Tolerance:** If a subscriber is unavailable or slow to process messages, it does not affect the publisher. The messages can be queued and delivered later.
- **Improved System Responsiveness:** The publisher can continue publishing messages without delay, even if the subscribers take time to process them.

Drawbacks:

- **No Immediate Feedback:** Since the publisher doesn't wait for acknowledgment, it doesn't immediately know whether the message was received or processed.
 - **Increased Complexity:** Asynchronous systems often require additional mechanisms for error handling, message retries, and dead-letter queues to deal with undelivered messages.
 - **Potential Message Ordering Issues:** In some cases, messages may be processed out of order, or with delays, making it harder to ensure strict sequencing.
-

Key Differences

Aspect	Synchronous Communication	Asynchronous Communication
Message Flow	Publisher waits for an acknowledgment from the subscriber	Publisher sends message and continues without waiting
Blocking/Non-blocking	Blocking (Publisher waits for a response)	Non-blocking (Publisher continues immediately)
Coupling	Tightly coupled (Direct dependency on subscriber response)	Loosely coupled (Publisher and subscriber are independent)
Response Time	Immediate feedback and acknowledgment	No immediate feedback; acknowledgment happens later
Scalability	Less scalable (Limited by subscriber response times)	Highly scalable (Can handle large volumes independently)
Latency	Higher latency (Depends on subscriber's processing time)	Low latency for publisher; processing can be deferred
Failure Handling	Failure directly impacts the publisher	More fault tolerant (Failures on subscriber side don't block the publisher)
Use Cases	Real-time, synchronous systems (e.g., RPC, request-response)	Event-driven, asynchronous systems (e.g., notifications, logging)
Complexity	Simpler, linear flow	More complex; requires error handling (retries, dead-letter queues)

Summary

- **Synchronous communication** in Pub/Sub is more suited for scenarios where the publisher needs immediate feedback from the subscriber, but it introduces higher latency, tighter coupling, and lower scalability. It is typically used in **real-time, transactional systems** where the result must be known immediately.
- **Asynchronous communication**, on the other hand, is highly scalable, resilient to failures, and non-blocking, making it ideal for **event-driven architectures** and systems where message processing can happen independently of the publisher's operations. However, it introduces more complexity in terms of handling failures, retries, and ensuring message delivery.

In most modern distributed systems, **asynchronous communication** is preferred for Pub/Sub because it enhances scalability, fault tolerance, and system performance.

7. What are the differences between push and pull delivery in Pub/Sub?

- **Push Delivery:** The Pub/Sub service sends messages to a subscriber's HTTP endpoint in real time as they arrive. The subscriber needs to have an HTTP server to accept incoming messages. Push subscribers must acknowledge messages after processing.
- **Pull Delivery:** The subscriber actively polls the Pub/Sub service to pull messages. Subscribers manually request and acknowledge messages. This approach provides more control over message flow and processing pace.

Use Cases:

- Push: Best for real-time message processing with low latency requirements.
- Pull: Better suited for subscribers with varying workloads or where you need more control over when and how messages are processed.

[Official Docs: Push and Pull](#)

	Pull	Push
Use case	Large volume of messages (many more than 1 per second).	Multiple topics that must be processed by the same webhook.
	Efficiency and throughput of message processing is critical.	App Engine Standard and Cloud Functions subscribers.
	Environments where a public HTTPS endpoint with a non-self-signed SSL certificate is not feasible to set up.	Environments where Google Cloud dependencies (such as credentials and the client library) are not feasible to set up.

- By default, Pub/Sub offers at-least-once delivery with no ordering guarantees on all subscription types.
- Alternatively, if messages have the same ordering key and are in the same region, you can enable message ordering.
- After the message ordering property is set, the Pub/Sub service delivers messages with the same ordering key in the order that the Pub/Sub service receives the messages.

8. How does Pub/Sub ensure at-least-once message delivery?

Pub/Sub ensures **at-least-once** delivery by requiring subscribers to acknowledge messages they have successfully processed. If a subscriber does not acknowledge a message within the specified **ack deadline**, Pub/Sub will reattempt delivery. Messages will continue to be redelivered until acknowledged, ensuring that they are received at least once.

However, this can lead to duplicate messages, so idempotent message processing is recommended.

[Official Docs: Message Delivery Guarantees](#)

9. What is the maximum size of a Pub/Sub message?

The maximum size of a single Pub/Sub message is **10 MB**. This includes both the message payload and any attributes associated with the message.

For large data, it is recommended to use a storage service like Google Cloud Storage and publish links to the data rather than the data itself.

[Official Docs: Quotas and Limits](#)

10. Explain the concept of message retention in Pub/Sub. What happens to unacknowledged messages?

Message retention defines how long Pub/Sub will retain a message after it has been published, even if it has been acknowledged. By default, Pub/Sub retains unacknowledged messages for **7 days**, though this can be configured up to **31 days**.

The following are the values for the **Message retention duration** option:

- Default value = 7 days
- Minimum value = 10 minutes
- Maximum value = 31 days

Unacknowledged messages remain available for redelivery until acknowledged or until the retention period expires, after which they are deleted.

[Official Docs: Message Retention](#)

11. How do **seek** and **snapshot** mechanisms work in Pub/Sub systems, and what are the advantages of using them for managing message delivery and processing in a subscriber application?

https://cloud.google.com/pubsub/docs/replay-overview#snapshot_overview

Seek and **Snapshot** mechanisms in Pub/Sub systems allow for more precise control over message consumption by subscribers, particularly in managing message delivery, recovery, and replay. Let's break down how they work and their benefits:

1. **Seek Mechanism**

Seek is a feature that allows a subscriber to **rewind or move forward** in the message stream to a specific point in time or message. Instead of processing messages in real-time, subscribers can "seek" to:

- **A specific timestamp:** The subscriber can seek to a specific time, effectively replaying messages that were published after that time.
- **A specific message ID:** The subscriber can seek to a specific message in the stream by its unique ID.
- **Start of the topic:** The subscriber can go back to the earliest available messages in the topic.
- **End of the topic:** The subscriber can skip forward to the most recent messages, ignoring the backlog.

Advantages of Seek:

- **Message Replay:** Subscribers can replay past messages, which is useful in scenarios where messages need to be reprocessed due to errors, system failures, or new features being deployed that require past data.
- **Backlog Recovery:** If a subscriber temporarily stops consuming messages and builds up a backlog, it can seek to a specific point to recover lost messages.
- **Debugging & Testing:** Developers can use seek functionality to debug issues by replaying the exact sequence of messages that triggered a problem.
- **Graceful Catch-up:** Allows for easier handling of subscriber restarts or failures by seeking to the point where it left off, ensuring no messages are missed.

2. Snapshot Mechanism

Snapshots in Pub/Sub refer to the ability to **capture the state of a subscription** at a specific point in time. A snapshot records the current message acknowledgment state for a subscription. This allows a subscriber to later **roll back** to this snapshot if needed, resuming message processing from that exact point.

How Snapshots Work:

- When a snapshot is created, Pub/Sub saves the acknowledgment state, which includes all the messages that have been processed (acknowledged) up to that point.
- If something goes wrong (e.g., an application bug or data corruption), the subscriber can revert to this snapshot to reprocess messages from the saved point onward.

Advantages of Snapshots:

- **Disaster Recovery:** In case of a failure or system crash, the subscriber can revert to a snapshot and resume processing from the saved state, avoiding data loss.
- **Fault Tolerance:** Snapshots provide a mechanism to reprocess a known good state, which helps ensure reliability when dealing with critical message processing.
- **Experimentation & Rollbacks:** If a subscriber has processed a batch of messages with incorrect logic, it can revert to a snapshot to "undo" the processing and reprocess the messages correctly.
- **State Persistence:** Snapshots capture the processing state over time, making it easier to maintain consistency even in distributed or fault-tolerant systems.

Use Case Scenario for Seek and Snapshots:

- **Data Pipelines:** If a data pipeline ingests messages but an error is detected in the processing logic, you can use the **Seek** function to replay messages from a certain point (after fixing the issue), ensuring that no data is lost.
- **System Recovery:** Suppose a failure occurs, leading to incomplete processing of messages. A **Snapshot** allows you to roll back to the last successful state and reprocess the data accurately, avoiding duplication or loss.

Key Differences Between Seek and Snapshot

Feature	Seek	Snapshot
Functionality	Moves subscriber to a specific point in time or message	Captures the current acknowledgment state for future rollback
Use Case	Replaying messages, recovering backlog	Rolling back to a previous state after errors or failures
Granularity	Can seek to any message ID or timestamp	Restores the subscriber to a precise acknowledgment state
Recovery Approach	More flexible for message reprocessing	Best for reverting to a known good state
State Management	No state is saved; simply repositions the subscriber	Saves the processing state at the time of the snapshot

Conclusion:

Both **Seek** and **Snapshot** mechanisms provide powerful tools for managing message processing in Pub/Sub systems, especially in terms of error handling, recovery, and replaying messages.

- **Seek** allows for granular control of message consumption, making it easier to recover or replay messages from any point in time.
- **Snapshots** offer a more robust rollback mechanism by capturing the acknowledgment state of the subscription, enabling the system to revert to a known good state after a failure.

12. How can you implement exactly-once processing using Pub/Sub?

Exactly-once processing can be achieved by implementing **idempotency** in your subscriber application. Since Pub/Sub guarantees at-least-once delivery but allows for the possibility of duplicates, your system should be able to handle duplicate messages gracefully.

Steps:

1. Assign unique identifiers (message IDs) to messages.
2. Use a database or cache to track processed message IDs and ensure that each message is processed only once.

[Official Docs: Message Deduplication](#)

13. What are the benefits of using Pub/Sub in a microservices architecture?

- **Decoupling**: Services can publish and subscribe to topics independently of each other.
- **Scalability**: Pub/Sub can scale to handle high volumes of messages and events.
- **Asynchronous Communication**: Pub/Sub allows microservices to communicate asynchronously, reducing dependencies and improving resiliency.
- **Global Reach**: Pub/Sub provides low-latency, reliable messaging across multiple regions, ideal for distributed systems.

[Official Docs: Microservices with Pub/Sub](#)

14. How does Pub/Sub handle message ordering?

Pub/Sub does not guarantee **message ordering** by default. However, **ordering keys** can be used to ensure that messages with the same key are delivered in the order they were published.

To enable message ordering:

1. Set the **ordering key** in each message.

2. Enable ordering on the Pub/Sub topic.

[Official Docs: Message Ordering](#)

15. Explain the concept of dead-letter topics in Pub/Sub.

A **dead-letter topic** is a special topic that collects messages that could not be successfully processed or acknowledged after a configurable number of delivery attempts (the **max delivery attempts**).

Once a message reaches the maximum delivery attempts, Pub/Sub moves it to the dead-letter topic, ensuring that the main subscription is not overloaded with undeliverable messages.

[Official Docs: Dead Letter Topics](#)

16. What are some common use cases for Google Cloud Pub/Sub?

- **Event-driven architectures**: Propagating state changes or events across microservices.
- **Real-time analytics**: Streaming data from IoT devices, social media, or other sources.
- **Data ingestion pipelines**: Sending data to systems like BigQuery, Dataflow, or Cloud Storage.
- **Decoupling systems**: Allowing systems to communicate without being tightly coupled.

Google Cloud Pub/Sub is a fully managed, scalable messaging service designed to integrate distributed systems and enable event-driven architectures. Here are some common use cases where Google Cloud Pub/Sub is widely applied:

1. Real-Time Analytics and Data Streaming

- **Use Case**: Collecting and processing large volumes of data from various sources, such as IoT devices, application logs, or clickstreams, in real-time.
- **Example**: Streaming user interaction data from a website or app to Google Cloud Dataflow or BigQuery for real-time analytics, dashboard updates, and business insights.

2. Microservices Communication

- **Use Case**: Enabling **asynchronous communication** between microservices, decoupling them to ensure scalability, reliability, and fault tolerance.

- **Example:** In a microservices-based e-commerce platform, Pub/Sub can be used to pass order information from the checkout service to the inventory, billing, and shipping services without blocking operations.

3. Event-Driven Architectures

- **Use Case:** Implementing an event-driven system where various services and applications respond to specific events by consuming messages from topics.
- **Example:** An online retailer can use Pub/Sub to trigger inventory updates, send notifications to customers, and start fulfillment processes when an order is placed.

4. IoT Data Ingestion

- **Use Case:** Handling streams of data from **IoT devices** (sensors, smart devices) at scale, distributing it to multiple downstream services for processing and analysis.
- **Example:** An IoT-based smart city project uses Pub/Sub to gather sensor data (e.g., traffic, temperature, pollution levels) and send it to data pipelines for aggregation, analysis, and visualization.

5. Log Aggregation and Monitoring

- **Use Case:** Collecting and distributing logs from various applications and systems in a centralized and scalable way.
- **Example:** A company collects logs from multiple application instances and sends them to Pub/Sub, which forwards the logs to Cloud Logging, Cloud Bigtable, or an external log processing system for monitoring and alerting.

6. Background Processing and Task Queuing

- **Use Case:** Offloading resource-intensive tasks from real-time operations to be processed asynchronously in the background.
- **Example:** A video processing application uses Pub/Sub to queue tasks such as video transcoding, where messages containing job details are sent to background workers for processing.

7. Multi-Cloud or Hybrid Cloud Integration

- **Use Case:** Enabling **communication between services** running across different cloud environments or between on-premise infrastructure and Google Cloud.
- **Example:** An organization using a hybrid architecture can use Pub/Sub to integrate on-premise systems (such as a data center) with cloud services for data synchronization or message exchange.

8. Distributing Workloads to Multiple Consumers

- **Use Case:** Load balancing messages across multiple worker nodes or instances in a distributed system, ensuring efficient resource utilization.
- **Example:** A financial institution uses Pub/Sub to distribute stock trading data across multiple consumers, each handling different types of processing, such as risk analysis or portfolio updates.

9. Notification Systems

- **Use Case:** Implementing push-based notification systems to inform users or systems of events in real-time.
- **Example:** A social media platform uses Pub/Sub to send real-time notifications to users about likes, comments, or messages, triggering immediate user interactions or updates.

10. Reliable Data Pipeline with Acknowledgments

- **Use Case:** Ensuring reliable and durable message delivery in a **data pipeline** with acknowledgment mechanisms and automatic retries.
- **Example:** A large-scale retailer sends transaction data to Pub/Sub, which ensures that the message is acknowledged and processed correctly by multiple downstream systems (e.g., CRM, inventory management) without message loss.

11. Decoupling Monolithic Applications

- **Use Case:** Breaking down a **monolithic application** into loosely coupled services that communicate via Pub/Sub, allowing the application to scale and evolve independently.
- **Example:** A legacy banking system can decouple its core components (such as payment processing, fraud detection, and customer notifications) using Pub/Sub to facilitate more flexible updates and scaling.

12. Data Replication and Synchronization

- **Use Case:** Replicating or synchronizing data between different databases, regions, or services, ensuring consistency across multiple environments.
- **Example:** A global e-commerce platform uses Pub/Sub to replicate user profile updates across multiple data centers, ensuring that changes are reflected consistently in all regions.

17. How does Pub/Sub handle message deduplication?

Pub/Sub does not natively guarantee exactly-once delivery but can achieve **deduplication** by leveraging **message IDs**. Subscribers can keep track of processed message IDs and discard any duplicates.

For exactly-once processing, applications need to be **idempotent** (able to handle duplicate messages without side effects).

Google Cloud Pub/Sub handles message deduplication to ensure that subscribers process each message only once, even in cases where duplicates are sent due to network issues, retries, or other factors. Pub/Sub is designed to be an **at-least-once delivery** system, which means a message may be delivered more than once, but deduplication mechanisms help mitigate this.

Here's how Pub/Sub handles message deduplication:

1. Message IDs for Uniqueness

Each message published to a Pub/Sub topic is assigned a **unique message ID** by the system. This message ID serves as an identifier that can be used by the subscriber to detect duplicate messages.

How It Works:

- When a subscriber receives a message, it can check the message ID.
- If the subscriber has already processed a message with that ID, it can safely discard the duplicate.

This approach helps clients ensure **idempotency** in message processing, where processing the same message more than once does not lead to unintended side effects.

2. At-Least-Once Delivery and Retries

Pub/Sub guarantees at-least-once delivery, which means that it may redeliver the same message under certain conditions:

- **Ack Deadline Expiry:** If a subscriber does not acknowledge (ACK) the message within the configured acknowledgment deadline, Pub/Sub will assume the message was not processed and will redeliver it.
- **Network Failures:** In cases of network interruptions, the system may not receive an acknowledgment from the subscriber, leading to retries.

To handle this, subscribers need to ensure they can handle repeated deliveries of the same message. This is where deduplication at the subscriber level becomes important.

3. Idempotent Processing

Since Pub/Sub does not guarantee exactly-once delivery, it is recommended that subscriber applications are designed to be **idempotent**. This means the processing logic should ensure that even if the same message is processed multiple times, the outcome remains consistent.

Deduplication can be achieved by:

- Using the **message ID** as a key to track already processed messages in a database or in-memory cache.
- Implementing logic in the subscriber to check whether a message has already been processed.

Example :

If a message instructs a system to charge a user's credit card, the subscriber can check a transaction log (with the message ID) to determine if the charge has already been processed. If it has, the subscriber can skip processing to avoid double charging the user.

4. Message Ordering and Deduplication

For subscriptions with **message ordering enabled**, Pub/Sub delivers messages with the same ordering key in the correct order. This ordering can also help subscribers detect duplicates, especially when coupled with idempotent processing techniques.

5. Best Practices for Deduplication

- **Use Unique Message IDs:** Always check the message ID before processing to avoid handling the same message multiple times.
- **Maintain a State Store:** Keep a record of processed message IDs in an external store like a database, key-value store, or in-memory cache. This is particularly important when processing is critical (e.g., financial transactions).
- **Design Idempotent Subscribers:** Ensure your message processing logic is idempotent, so the system behaves correctly even if it processes the same message multiple times.
- **Set Appropriate Ack Deadlines:** Choose the right acknowledgment deadlines based on how long your processing takes. This reduces the likelihood of messages being redelivered due to missed acknowledgments.

Conclusion:

While Google Cloud Pub/Sub uses **message IDs** to help with deduplication, subscribers are responsible for ensuring that they process messages idempotently. Pub/Sub provides at-least-once delivery guarantees, which means it may occasionally deliver duplicate messages, especially if acknowledgments are missed. Handling deduplication and ensuring idempotent processing on the subscriber side is crucial for ensuring system integrity when using Pub/Sub.

[Official Docs: Idempotency and Deduplication](#)

18. What is the purpose of the ack deadline in Pub/Sub, and how can you modify it?

The **ack deadline** is the maximum time a subscriber has to acknowledge a message after receiving it. By default, it is **10 seconds**, but can be extended up to **600 seconds**.

Subscribers can modify the ack deadline for long-running message processing tasks using the `modifyAckDeadline` method, allowing more time to process a message before Pub/Sub attempts redelivery.

[Official Docs: Ack Deadlines](#)

19. How does Pub/Sub integrate with other Google Cloud services?

- **Dataflow**: Streaming data processing using Pub/Sub as a source or sink.
- **BigQuery**: Directly ingest Pub/Sub data into BigQuery for analysis.
- **Cloud Functions**: Trigger functions when messages are published to a topic.
- **Cloud Storage**: Store large payloads and send notifications via Pub/Sub when new files are added.

[Official Docs: Pub/Sub Integration](#)

20. What are some best practices for designing topics and subscriptions in Pub/Sub?

- Use **separate topics** for different event types or services.
- Avoid creating too many topics or subscriptions to reduce overhead.
- Enable **ordering keys** for messages that require ordered delivery.

- Use **dead-letter topics** to handle unprocessable messages.
- Use **message filtering** to avoid delivering unnecessary messages to subscribers.

[Official Docs: Best Practices](#)

21. How can you implement message filtering in Pub/Sub?

Message filtering allows subscribers to receive only messages that match certain criteria, reducing unnecessary message processing. Filters are applied using **filter expressions** based on message attributes

Example:

```
attributes.domain = "finance"
```

language-sql

This filter would ensure that only messages with the attribute `domain = "finance"` are delivered to the subscriber.

[Official Docs: Message Filtering](#)

22. Explain the concept of message retention and how it affects your Pub/Sub usage.

Message retention defines how long Pub/Sub keeps messages available for redelivery after being published. Retention is especially important for cases where subscribers may need to catch up after being offline. Retention affects costs, as longer retention times may result in higher storage costs.

By default, unacknowledged messages are retained for 7 days, but you can configure retention up to 31 days.

[Official Docs: Message Retention](#)

23. What security features does Pub/Sub offer to protect your data?

- **IAM (Identity and Access Management)**: Fine-grained access control to topics and subscriptions, specifying who can publish, subscribe, or manage resources.

- **Encryption:** Pub/Sub encrypts data at rest and in transit using Google-managed keys. For additional security, customers can use **Customer-Managed Encryption Keys (CMEK)**.
- **Private Google Access:** Pub/Sub can be accessed over private IP addresses using **VPC Service Controls** to prevent public exposure.

[Official Docs: Security](#)

24. How would you monitor and troubleshoot a Pub/Sub system in production?

To monitor and troubleshoot a Pub/Sub system, you can:

1. **Cloud Monitoring:** Use Google Cloud Monitoring to track message metrics (e.g., publish rate, subscription backlog, ack rates).
2. **Cloud Logging:** View logs related to published messages, subscriptions, and message delivery failures.
3. **Alerts:** Set up custom alerts for conditions such as large backlogs or message delivery failures.
4. **Dead-letter Topics:** Use dead-letter topics to investigate undelivered messages.

[Official Docs: Monitoring](#)